

Graph Database Management System

Che Liu
A53209595

Overview

Overview

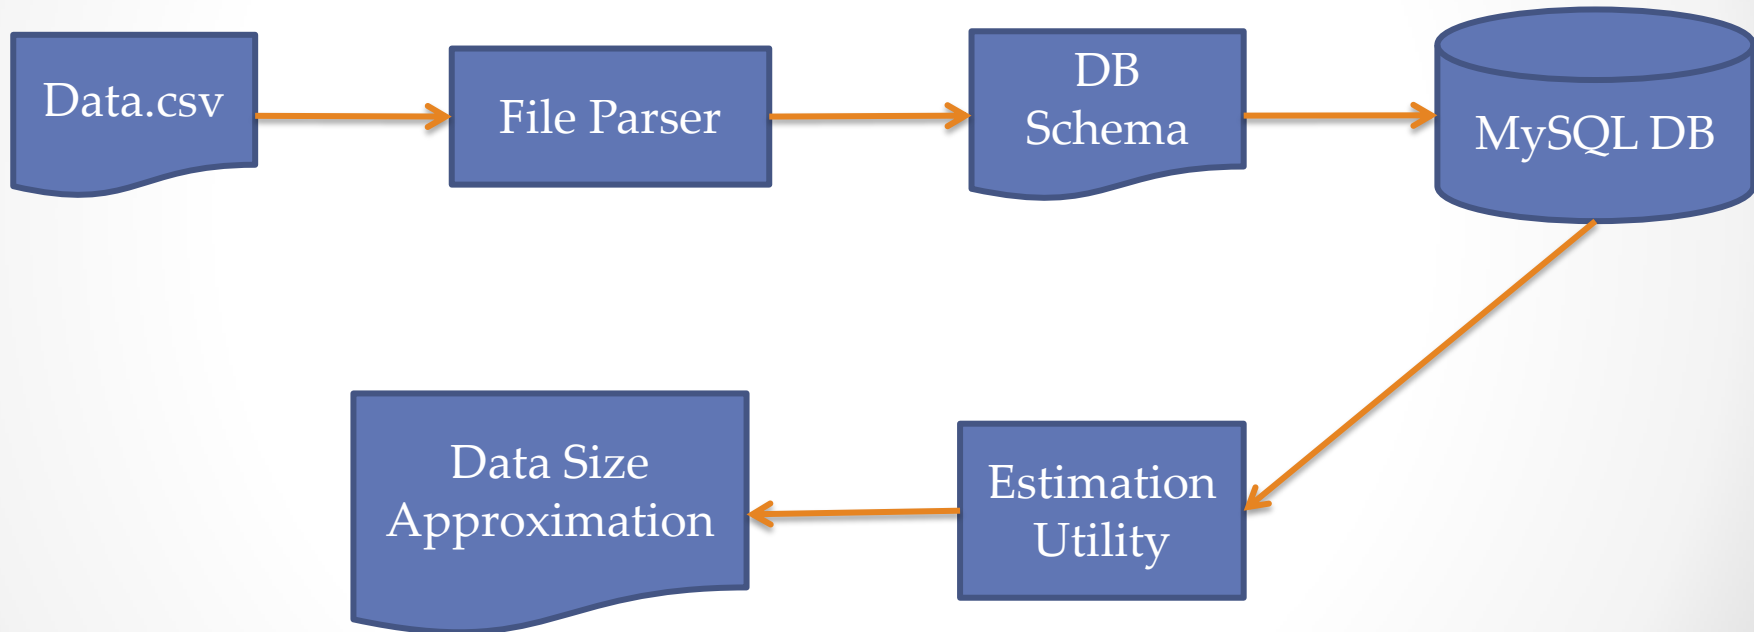
- **File Parsing**

- Generate MySQL Schema
- Import data into database
- Data estimation

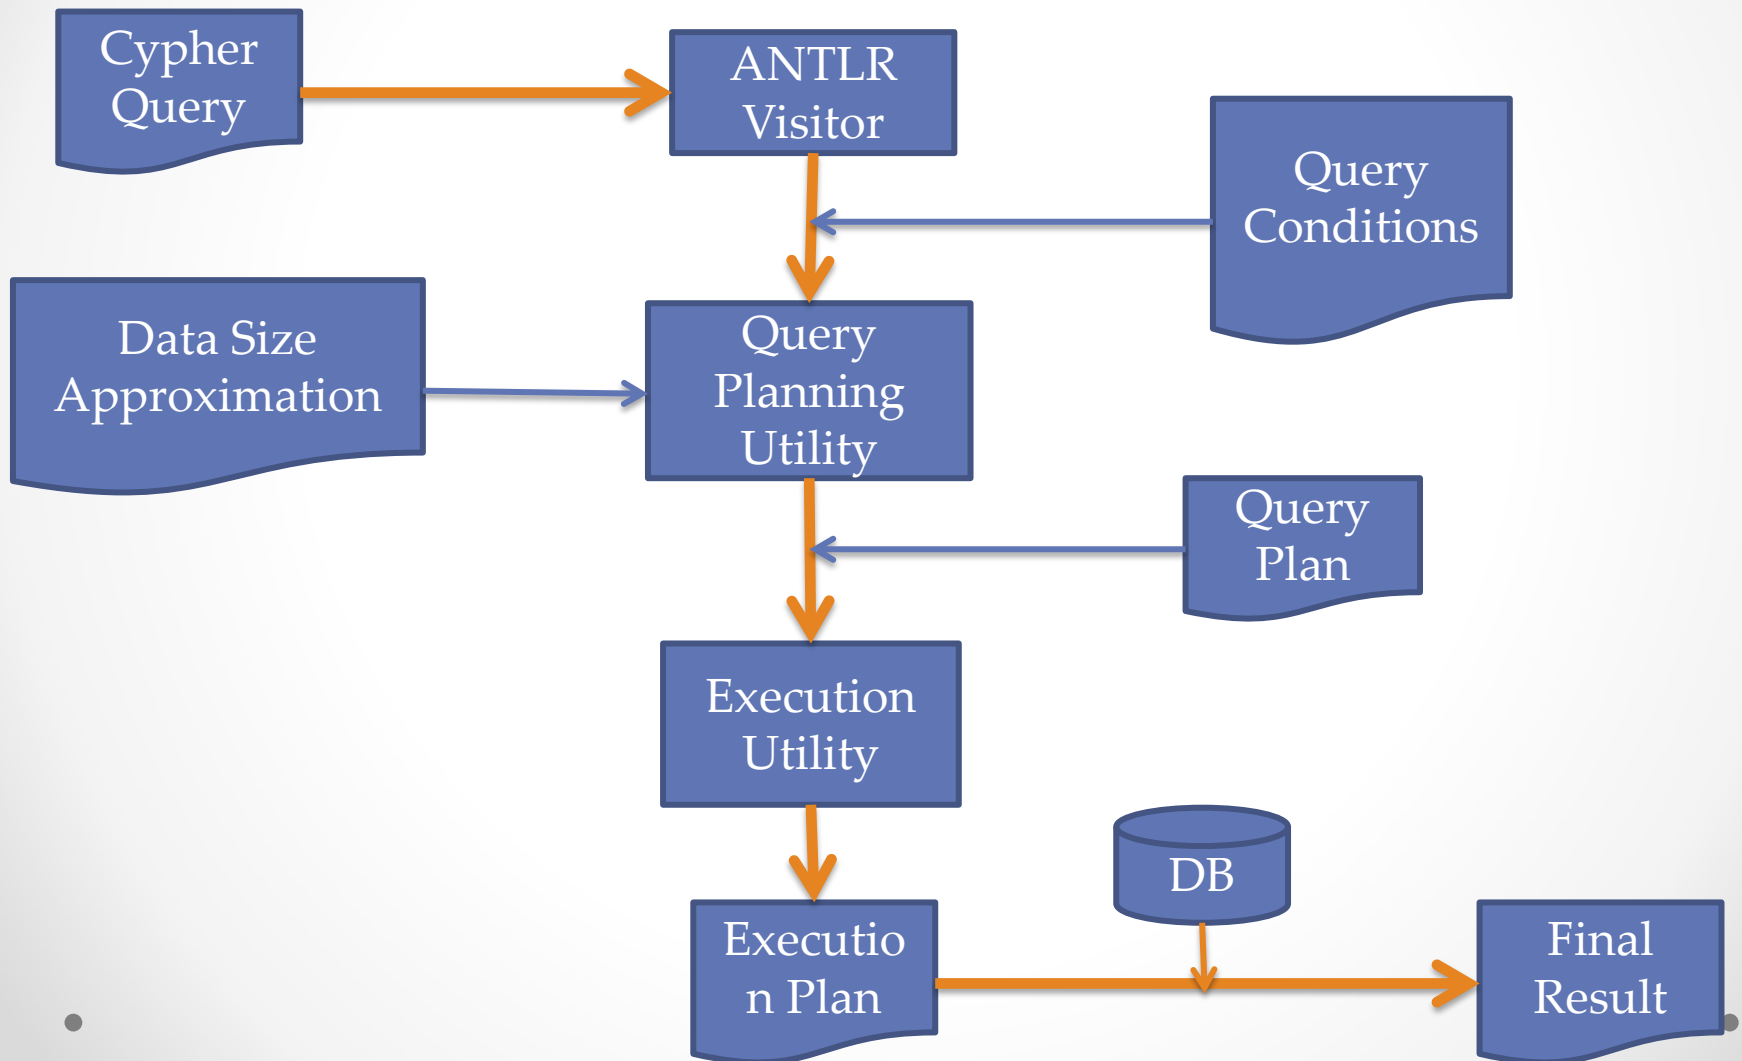
- **Query Execution**

- Query parsing
- Generate query plan, execution plan
- Produce results

File Parsing



Query Execution



System Components

- File import utility
- MySQL Database
- Parsing query string
- Data size approximation
- Query planning and size estimation
- Query execution

File import utility

Parsing

- Buffered Reader to read file
- Convert to JSON Object

SQL

- Generate SQL Schema
- Generate insert statements

Import DB

- JDBC
- Batch insert to database

File Import Utility

File import utility

- **First scan:**
 - Create metadata:
 - Collect schema of different type of nodes
 - Convert type to SQL types.
 - Set up data tables and indices.
- **Second scan:**
 - Collect nodes and insert them to database
 - Assume ID in metadata is unique among nodes.
 - Assign each node with a unique GID (Global ID).
 - Insert node and its properties to DB.

MySQL Database Design

MySQL Database

- **Column-based storage:**
 - Same properties of all the nodes are stored in the same table.
 - GID – value (key-value) storage.
 - Index on all primary keys
 - Index on frequent used values: TID, UID, Label, Name...
- **Small Edge table**
 - Unique index number
 - Node1 – relation_type – Node2
 - Additional information (nodeLabel) for fast lookup
 - If small enough, it can stay in memory.

Database Graph

| P_Property1 |
|-------------|
| gid |
| value |

| NodeLabel |
|-----------|
| gid |
| label |

| P_Property2 |
|-------------|
| gid |
| value |

| Edge |
|--------------|
| eid |
| node1 |
| node2 |
| rel_type |
| relationship |
| node1Type |
| node2Type |

| ObjectType |
|------------|
| gid |
| type |

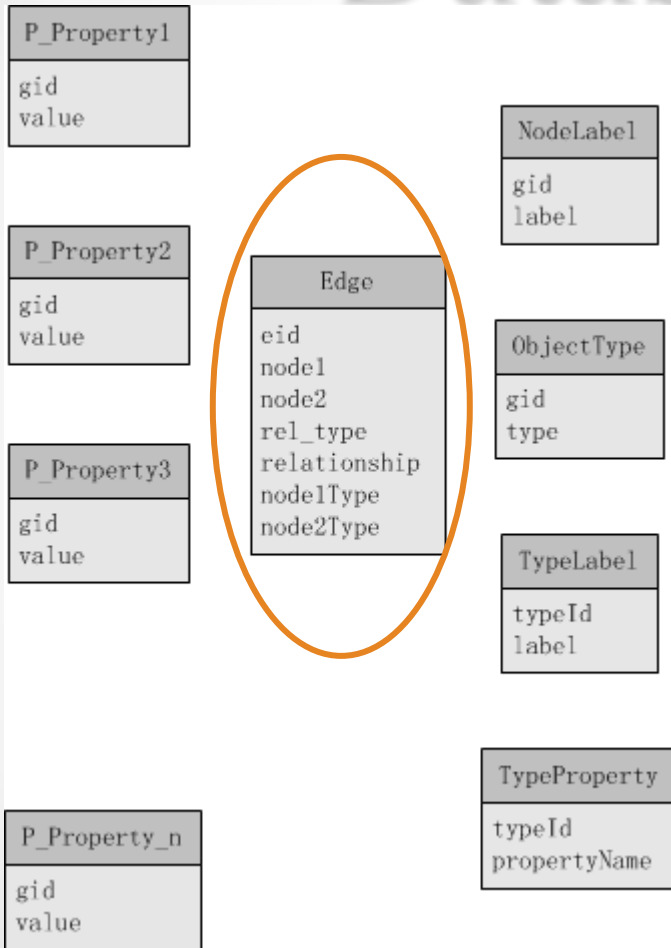
| P_Property3 |
|-------------|
| gid |
| value |

| TypeLabel |
|-----------|
| typeId |
| label |

| TypeProperty |
|--------------|
| typeId |
| propertyName |

| P_Property_n |
|--------------|
| gid |
| value |

Database Graph

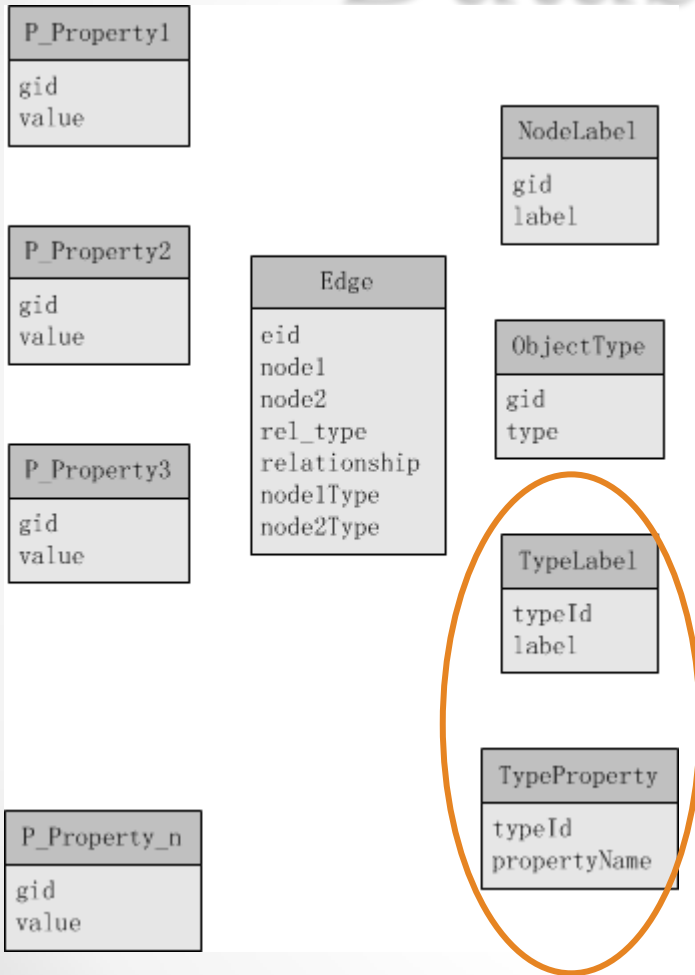


Edge Table

- Small enough to be kept in memory
- Node – Edge – Node
- Keep node type in table to accelerate lookup
- Everything have index.

Database Graph

Database Meta

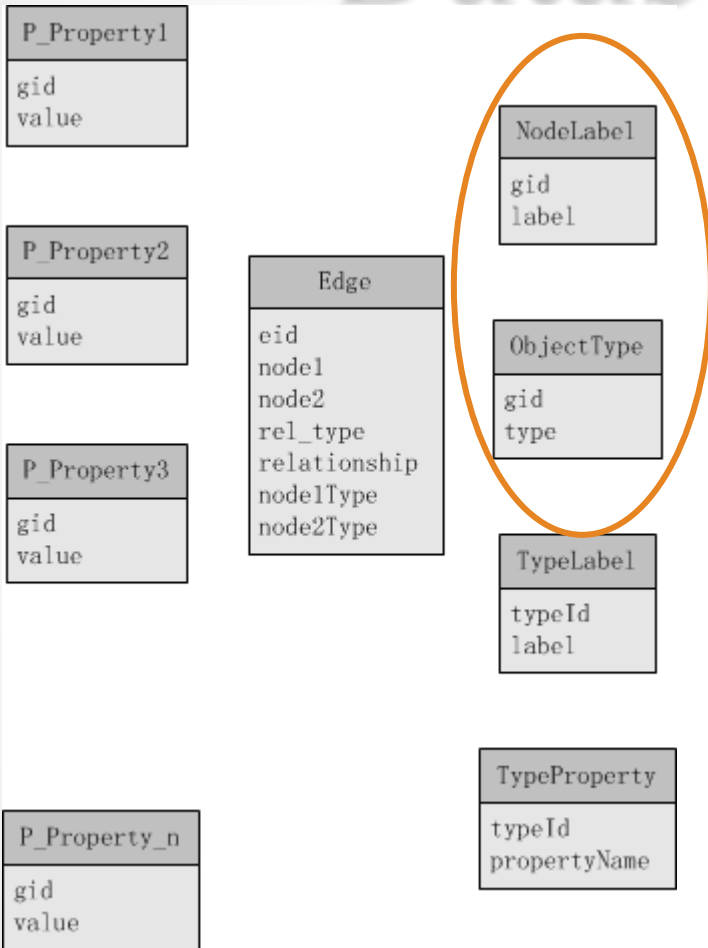


- Assumption:
 - Node with same label have the same schema.
 - One node can have many labels
- Assign each kind of node a type number
 - {Tweet} => 1
 - {Hashtag, LongHashtag} => 2
 - {Hashtag} => 3
- Keep schema in metadata

Database Graph

Node Meta

- Gid – value pairs.
- Information about labels and type of a node.



Database Graph

Properties table

- Gid – value pairs.
- Contents in the node

| P_Property1 |
|-------------|
| gid |
| value |

| P_Property2 |
|-------------|
| gid |
| value |

| P_Property3 |
|-------------|
| gid |
| value |

| P_Property_n |
|--------------|
| gid |
| value |

| Edge |
|--------------|
| eid |
| node1 |
| node2 |
| rel_type |
| relationship |
| node1Type |
| node2Type |

| NodeLabel |
|-----------|
| gid |
| label |

| ObjectType |
|------------|
| gid |
| type |

| TypeLabel |
|-----------|
| typeId |
| label |

| TypeProperty |
|--------------|
| typeId |
| propertyName |

Node lookup

| P_Property1 |
|--------------|
| gid value |

| P_Property2 |
|--------------|
| gid value |

| P_Property3 |
|--------------|
| gid value |

| P_Property_n |
|--------------|
| gid value |

| Edge |
|-----------------------------------------------------------------------------|
| eid node1 node2 rel_type relationship node1Type node2Type |

| NodeLabel |
|--------------|
| gid label |

| ObjectType |
|-------------|
| gid type |

| TypeLabel |
|-----------------|
| typeId label |

| TypeProperty |
|------------------------|
| typeId propertyName |

- GID

GID

Node lookup

| P_Property1 |
|--------------|
| gid value |

| P_Property2 |
|--------------|
| gid value |

| P_Property3 |
|--------------|
| gid value |

| P_Property_n |
|--------------|
| gid value |

| Edge |
|-----------------------------------------------------------------------------|
| eid node1 node2 rel_type relationship node1Type node2Type |

| NodeLabel |
|--------------|
| gid label |

| ObjectType |
|-------------|
| gid type |

| TypeLabel |
|-----------------|
| typeId label |

| TypeProperty |
|------------------------|
| typeId propertyName |

GID

type

- GID
- Get Node Type

Node lookup

| P_Property1 |
|--------------|
| gid value |

| P_Property2 |
|--------------|
| gid value |

| P_Property3 |
|--------------|
| gid value |

| P_Property_n |
|--------------|
| gid value |

| Edge |
|-----------------------------------------------------------------------------|
| eid node1 node2 rel_type relationship node1Type node2Type |

| NodeLabel |
|--------------|
| gid label |

| ObjectType |
|-------------|
| gid type |

| TypeLabel |
|-----------------|
| typeId label |

| TypeProperty |
|------------------------|
| typeId propertyName |

GID

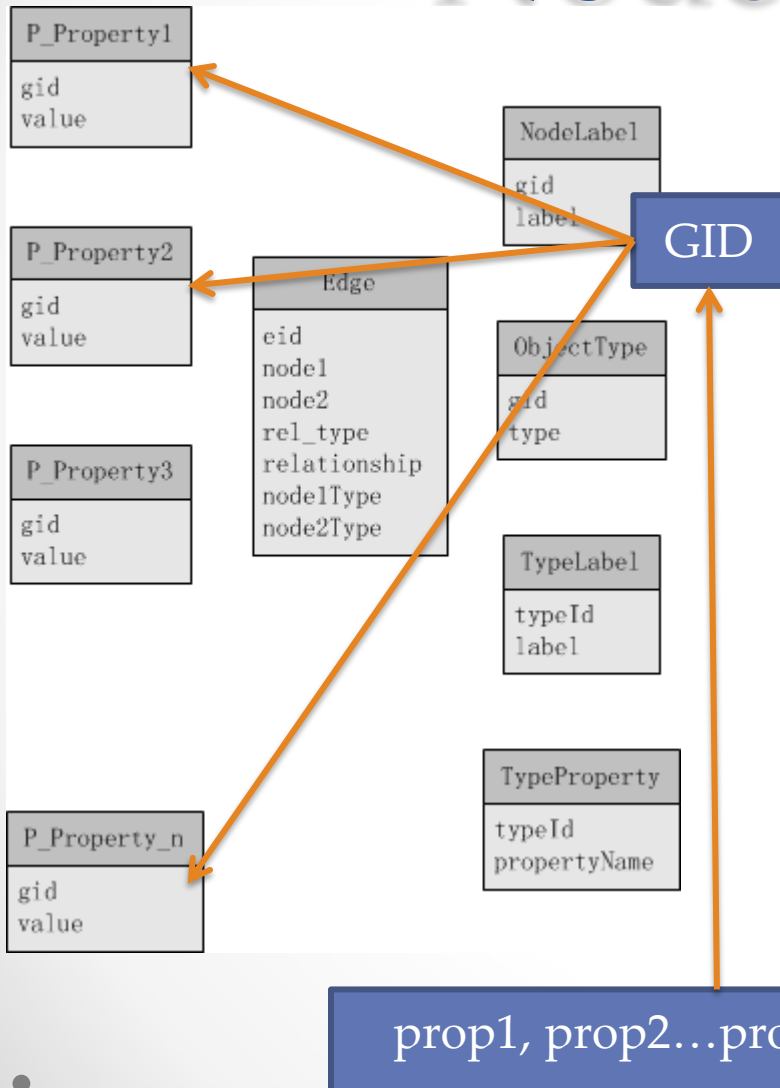
type

prop1, prop2...prop_n

- GID
- Get Node Type

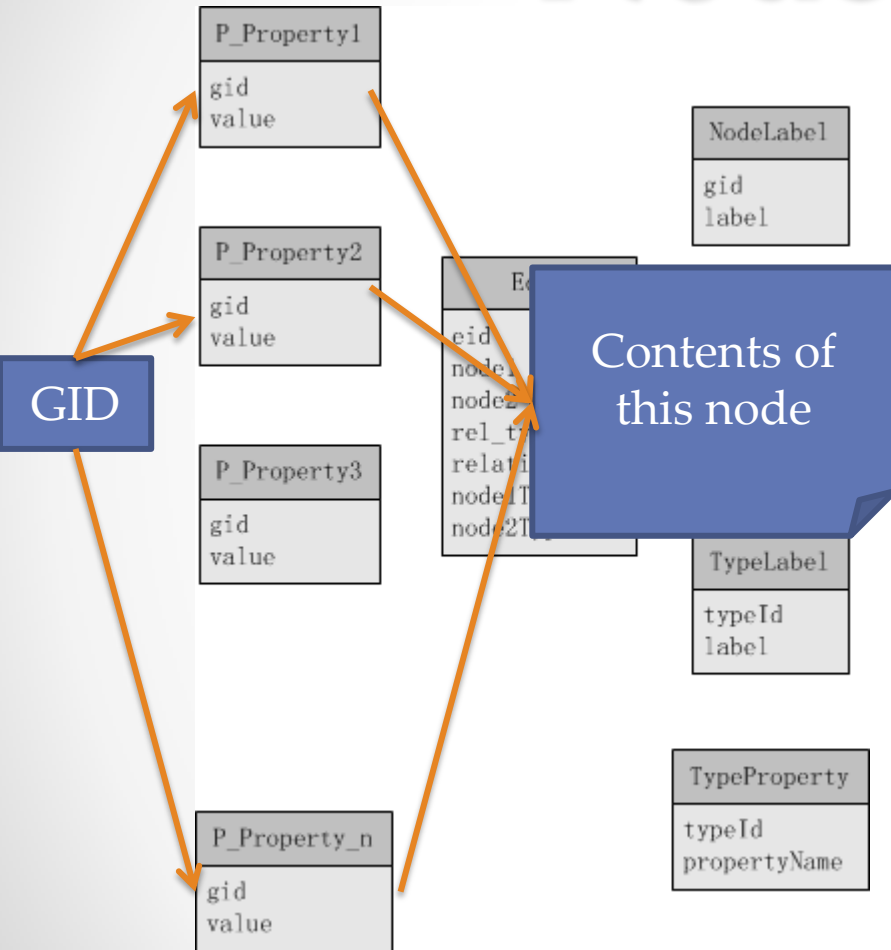
Get property names of this type of node

Node lookup



- GID
- Get Node Type
- Get property names of this type of node
- Get contents by property name and GID

Node lookup



- GID
- Get Node Type
- Get property names of this type of node
- Get contents by property name and GID
- Combine these results

Data Size Estimation

Data Size Estimation

- Run when importing finish.
- Provide rough insight of current dataset.
- Once created, never changes.
- Useful for size and cost estimation.

Data Size Estimation

| Name | Description |
|--------------------------------------------|-------------------------------------------------------------------------------|
| NumberOfNodes | Total number of nodes |
| NumberOfRelations | Total number of relations |
| MaxNodeIn | Maximum in-degree of node |
| MaxNodeOut | Maximum out-degree of node |
| LabelRelation | Number of relations with each label(rel_type) |
| LabelNodes | Number of nodes with same label set |
| PropCntOfNode | Number of nodes with different value on each property |
| NodeLabelIncoming/ NodeLabelOuting | Number of incoming/outgoing edges from nodes with same label |
| NodeRelationInEdge/ NodeRelationOutEdge | Number of incoming/outgoing edges of the same type from nodes with same label |

Query Parsing

Query Parsing

- Visitor Method to walk the parse tree
- **Constants** are evaluated immediately
- **Variables** and related **expressions** are collected together.
- **Return** expressions are parsed and collected into somewhere else.
- Patterns are extracted after visiting **match** clause

Query Parsing Example

Query:

Match (a) – [r1:hashtagUsedIn] → (b:tweet:longtweet) ,
(c:user {id = 1234}) – [r2:mentioned*1..3] – (b),

MATCH (e) <- [r3] - (f)

Where a.name = d.name AND b = f

return a.id, r1, b, f

Query Parsing Example

Query:

Match (a) – [r1:hashtagUsedIn] → (b:tweet:longtweet) ,
(c:user {id = 1234}) – [r2:mentioned*1..3] – (b),

MATCH (e) <- [r3] - (f)

Where a.name = d.name **AND** b = f **AND** e.id = 1000

return a.id, r1, b, f

Query Parsing Example

- Collect conditions on each variable:

- a: -
- r1 :hashtagUsedIn
- b :tweet:longtweet
- c :user id = 1234
- r2 :mentioned*1..3
- e .id = 1000
- r3: -
- d: -

MATCH

(a) – [r1:hashtagUsedIn] →
(b:tweet:longtweet) ,
(c:user {id = 1234}) –
[r2:mentioned*1..3] – (b),

MATCH

(e) <- [r3] - (d)

WHERE

a.name = d.name AND
b = d AND
e.id = 1000

return a.id, r1, b, d

Query Parsing Example

- Collect conditions between variables:

- `a.name = d.name`
- `b = d`
- `r1 != r2`

MATCH

(a) – [r1:hashtagUsedIn] →
(b:tweet:longtweet) ,
(c:user {id = 1234}) –
[r2:mentioned*1..3] – (b),

MATCH

(e) <- [r3] - (d)

WHERE

`a.name = d.name` **AND**

`b = d` **AND**

`e.id = 1000`

return `a.id`, r1, b, d

Query Parsing Example

- Collect query patterns:
- a-r1->b
- c-r2->b
- e<-r3-d
- Long patterns are also supported.

MATCH

(a) – [r1:hashtagUsedIn] ->
(b:tweet:longtweet) ,
(c:user {id = 1234}) –
[r2:mentioned*1..3] – (b),

MATCH

(e) <- [r3] - (d)

WHERE

a.name = d.name **AND**

b = d **AND**

e.id = 1000

return a.id, r1, b, d

Query Parsing Example

- Collect return variables:

- a.id
- b
- d
- r1

MATCH

(a) – [r1:hashtagUsedIn] →
(b:tweet:longtweet) ,
(c:user {id = 1234}) –
[r2:mentioned*1..3] – (b),

MATCH

(e) <- [r3] - (d)

WHERE

a.name = d.name **AND**

b = d **AND**

e.id = 1000

return a.id, r1, b, d

Query Parsing

- Parse Results:
 - A list of matching patterns (node – edge - node)
 - A map containing all the conditions on each individual variable
 - A list containing all equations between variables
 - A list containing what result to return

Query Planning

Query planning

- Parse Results
- Greedy operator ordering algorithm
- Size estimation on each plan
 - Data Size Estimation
 - Size estimation of previous plan
- Produce plan tree

Query planning

- **Parse Results:**
 - A list of matching patterns (**PathList**)
 - A map containing all the conditions on each individual variable (**VarCondition**)
 - A list containing all equations between variables (**Equality**)
 - A list containing what result to return (**ReturnList**)

Greedy Operator Ordering

Input: Query graph Q
Result: Logical plan P that covers all nodes from Q

```
1  $\mathcal{P} \leftarrow []$  ▷ PlanTable
2 foreach  $n \in Q$  do ▷ every node in the query graph
3    $T \leftarrow \text{constructLeafPlan}(n)$  ▷ take selections into account
4    $\mathcal{P}.\text{insert}(T)$ 
5 do
6    $\text{Cand} \leftarrow []$  ▷ candidate solutions
7   foreach  $P_1 \in \mathcal{P}$  do
8     foreach  $P_2 \in \mathcal{P}$  do
9       if  $\text{CanJoin}(P_1, P_2)$  then
10          $T \leftarrow \text{constructJoin}(P_1, P_2)$ 
11          $\text{Cand}.\text{insert}(T)$ 
12   foreach  $P_1 \in \mathcal{P}$  do
13      $T \leftarrow \text{constructExpand}(P_1)$ 
14      $\text{Cand}.\text{insert}(T)$ 
15   if  $\text{Cand}.\text{size} \geq 1$  then
16      $T_{\text{best}} \leftarrow \text{pickBest}(\text{Cand})$  ▷ pick the plan with the smallest cost
17     foreach  $T \in \mathcal{P}$  do
18       if  $\text{covers}(T_{\text{best}}, T)$  then
19          $\mathcal{P}.\text{erase}(T)$  ▷ delete plans covered by  $T_{\text{best}}$ 
20      $T_{\text{best}} \leftarrow \text{applySelections}(T_{\text{best}})$ 
21      $\mathcal{P}.\text{insert}(T_{\text{best}})$ 
22 while  $\text{Cand}.\text{size} \geq 1$ 
23 return  $\mathcal{P}[0]$ 
```

Query planning

- **Basic rules**
 - Push down selections (**ConstructLeafPlan**)
 - Choose the minimum cost plan at each step(**pickbest**)
 - Reduce candidate size by removing plans covered by the best plan. (**covers**)

Greedy Operator Ordering

- **Plan**
 - Name of this plan
 - Variables that this plan apply to
 - Estimated cost of this plan
- **Size estimation of a plan**
 - Data size estimation
 - Plan table that this plan is applied to
 - Type of plan

Greedy Operator Ordering

- **PlanTable**

- Set of plans that are used
 - Represented by Query Plan Tree
- Variables that covered by this plan table
- Current estimated size
- Cost to execute these plans
 - Sum of cost of all plans in this plan tree

Greedy Operator Ordering

- **ConstructLeaf**
 - For each node, apply selections immediately
 - Serve as leaf node in the plan tree.
 - Selection related with variables are stored in **VarCondition**
- **Produced plans:**
 - ScanByIdPlan, ScanByPropertyPlan, ScanByLabelPlan, AllNodeScanPlan
 - FilterConstraintsPlan (not a leaf plan)

Greedy Operator Ordering

- **ConstructJoin**

- If two plan tables have common variables, or if there are equalities between two tables, produce NodeHashJoinPlan
 - $[a, b, c] \text{ \& } [a, d, e]$
 - $[a, b] \text{ \& } [c, d] \text{ ON } b = c$
- Otherwise, use Cartesian Product

- **Produced plans:**

- CartesianProductPlan, NodeHashJoinPlan

Greedy Operator Ordering

- **ConstructExpand**
 - If some node in the PlanTable have an edge linked to it, produce ExpandAllPlan
 - Table: {a, r, c}, Edge: (a)-[r2]->(b)
 - If both end of the edge lies in the node set of a PlanTable, produce ExpandIntoPlan
 - Table: {a, b, r}, Edge: (a)-[r2]-(b)
 - If this expansion is variable length expand, construct RangedExpansion
- **Produced plans:**
 - ExpandAllPlan, ExpandIntoPlan
 - RangeExpandAllPlan, RangeExpandIntoPlan

Greedy Operator Ordering

- **AddAdditionalFilter**

- Expand operation add variables into PlanTable
- Check if there are additional constraints on these variables.
 - Table: {a, r, b}, ExpandAll: (a)->[r2]->(c)
 - Suppose additional condition:
 - $r \neq r2, c.id = 2$

- **Produced plans:**

- FilterConstraintsPlan, FilterRelationEqualityPlan

Greedy Operator Ordering

- **ProduceFinalResult**
 - Look into return variables and conditions
(ReturnList)
- **Produced plans:**
 - ProduceResultPlan

Greedy Operator Ordering

- **Types of generated plans:**
 - ScanByIdPlan, ScanByPropertyPlan, ScanByLabelPlan, AllNodeScanPlan
 - CartesianProductPlan, NodeHashJoinPlan
 - ExpandAllPlan, ExpandIntoPlan
 - RangeExpandAllPlan, RangeExpandIntoPlan
 - FilterConstraintsPlan, FilterRelationEqualityPlan
 - ProduceResultPlan

Query Plans

- **ScanByIDPlan(NodeVar)**
 - Id is unique
 - **Estimated size** = 1
 - **Estimated cost** = 1
 - **Variable** = Node

Query Plans

- **ScanByLabelsPlan(NodeVar, Labels[])**
 - **Variable** = Node
 - **Estimated size** = Minimum number of nodes in the dataset with one of the labels provided
 - Data size estimation provides:
LabelNodes -> Number of nodes with same label set
 - **Estimated cost** = estimated size
 - Don't need to scan whole dataset since labels are indexed.

Query Plans

- **ScanByPropertyPlan**
- **(NodeVar, Property)**
 - **Variable** = Node
 - **Estimated size** = $\text{numberOfNodes} / \text{number of distinct values of this property}$.
 - **Estimated cost** =
 - Estimated size, if this property is indexed
 - Number of nodes, if not indexed.

Query Plans

- **AllNodeScanPlan(NodeVar)**
 - **Variable** = Node
 - **Estimated size** = number of nodes in dataset
 - **Estimated cost** = estimated size

Query Plans

- **FilterByConstraintPlan(Table, Condition)**
 - Takes a table, filter its result by the condition provided, and return the filtered result
 - **Variable** = Node
 - **Params** = condition
 - **Estimated size** =
 - $\min(\text{number of records in previous table}, \text{maximum number of records in the dataset})$
 - **Estimated cost** = number of nodes in previous table

Query Plans

- **FilterByRelationEquality**
(Table, Equality)
 - Takes a table, filter its result by the relation equation provided, and return the filtered result
 - **Variable** = Variables in the equation
 - **Estimated size** =
 - Table size (Not easy to estimate)
 - **Estimated cost** =
 - number of records in previous table

Query Plans

- **ExpandAllPlan(Table, conditions, PatternEdge)**
 - Expand the result set by the edge provided
 - **Variable** = Start variable of the edge
 - Start variable is the node in this edge that also occurs in the provided table.
 - **Size estimation** is same as
 - Pattern $(a:X:Y)-[:T1|:T2]->(b:W:Z)$
 - a and b each has two labels
 - Upper bound on the cardinality: C
 - $T1 = \min(X-[:T1]->(), Y-[:T1]->()) + \min(X-[:T2]->(), Y-[:T2]->())$
 - $T1' = \min(()-[:T1]->W, ()-[:T1]->Z) + \min(()-[:T2]->W, ()-[:T2]->Z)$
 - $C = \min(T1, T1')$

Query Plans

- **ExpandAllPlan(Table, conditions, PatternEdge)**
 - If there are no conditions provided:
 - **Estimated size** =
 - $\text{tableSize} * \text{maxDegreeOfNode}$ (conservative estimation)
 - $\text{tableSize} * \text{NumberOfEdges} / \text{NumberOfNodes}$
 - (rough estimation)

Query Plans

- **ExpandIntoPlan(Table, conditions, PatternEdge)**
 - Similar to ExpandAllPlan
 - Only difference is:
 - **Cost estimation** = min(size estimation, number of records in previous table)
 - If previous table is small, only a scan through previous table is needed.

Query Plans

- **RangedExpandAllPlan(Table, conditions, PatternEdge, Range)**
 - Range: (min, max)
 - Similar to ExpandAllPlan
 - We have to perform ExpandAll operation for MAX times, so:
 - **Cost Estimation** =
 - (cost of similar ExpandAllPlan) \wedge MAX

Query Plans

- **RangedExpandIntoPlan(Table, conditions, PatternEdge, Range)**
 - Range: (min, max)
 - Similar to ExpandIntoPlan
 - We have to perform ExpandInto operation for MAX times, so:
 - **Cost Estimation** =
 - (cost of similar ExpandIntoPlan) \wedge MAX

Query Plans

- **CartesianProductPlan(Table1, Table2)**
 - Perform CartesianProduct on two tables
 - No other condition needed.
 - **Variable**: All the variables in both tables
 - **Size estimation** =
 - Size of table1 * size of table2
 - **Cost Estimation** = size estimation

Query Plans

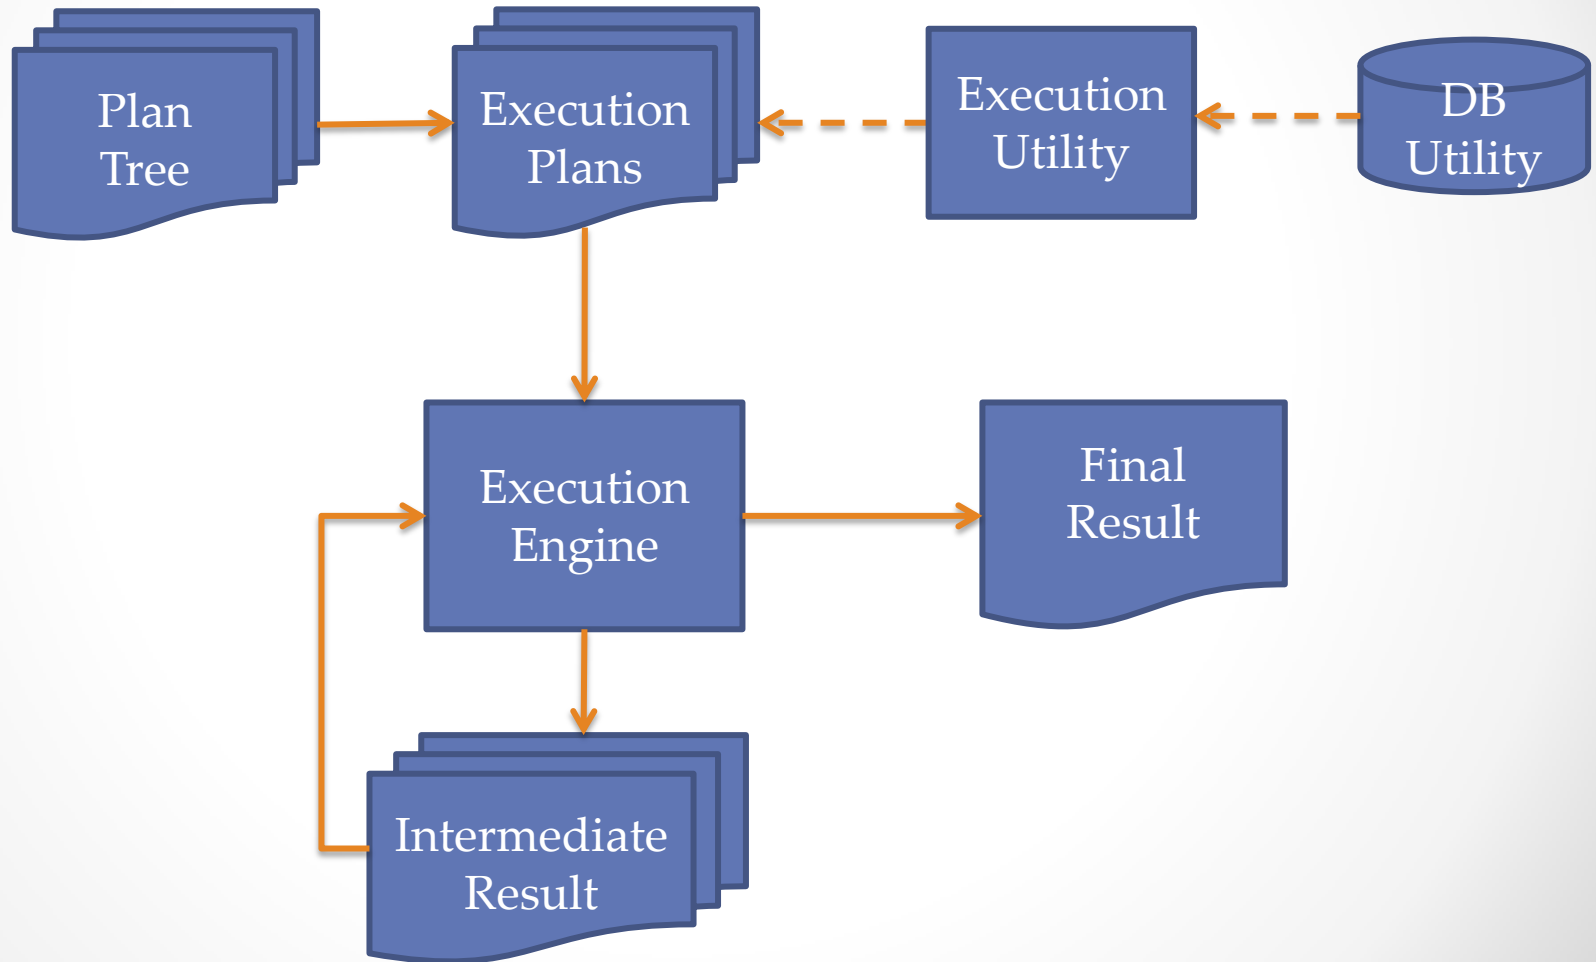
- **NodeHashJoinPlan(Table1, Table2, VarEqualities)**
 - Perform Hash Join on two tables, hash is created by the common variables or the equation provided.
 - **Variable:** All the variables in both tables
 - **Size estimation** =
 - Minimum size of two tables
 - **Cost Estimation** =
 - Sum of size of two tables

Query Execution

Query execution

- **Best possible Plan Tree**
- **Convert query plan into execution plan**
 - Done in the execution plan itself.
 - Each query plan have a counterpart execution plan
- **Execution Plan takes intermediate result and produces intermediate result (ResultTable)**
- **Execution Engine takes care of execution plans**
- **Produce final result and the SQL statements used**

Query execution



Query Plan Tree

- Query Plans is a **plan tree**.
- **Post-order traversal** to turn the tree to a list
- Easy to implement on a single-threaded process

Execution Engine

- **Stack-based Post order traversal**
 - Get the next execution plan, and operands needed. Different execution plans takes different number of operands (tables)
 - Fetch operand from stack
 - Call execution plan,
 - Put result back to stack

Execution Utility

- Contains methods to convert query into SQL statement
- Eg:
 - GetNodeBy: Get GID of a node by (its label, property and value...)
 - GetEdgeBy: Get Edge ID by the start/end/relation type/...
- Multiple conditions can be combined together to form a query.
 - GetNodeBy can be assigned both node label and node property as query condition at the same time
- Decouple data representation from data model.

Database Utility

- Execute the SQL query, and parse result into different types of Java Objects.
 - `getAsList`, `getAsMap`, `getAsTable`, `getAsString...`
- Record the SQL statements used.

Execution Plans

- **AllNodeScanExec(Plan)**
 - Get GIDs of all nodes and put them into result table. Returns a new result table.
 - Number of nodes is >2M, takes about 10s to finish, and consumes about 16M space
 - ResultTable is **huge**, not good for future plans

Execution Plans

- **ScanByIdExec(Plan)**
 - Get GIDs of nodes with some ID. Returns a new result table.
 - Only returns 1 result. Lookup is fast since ID is indexed in backend database.
 - **SQL: SELECT gid FROM P_id where value = SOME_ID**

Execution Plans

- **ScanByLabelExec(Plan)**
 - Get GIDs of all nodes with the label occurred in that plan. Returns a new result table.
 - Average number of nodes with same label is not small
 - Reduced intermediate result size.
 - Labels of all nodes are stored in **NodeLabel** table in MySQL.

Execution Plans

- **ScanByPropertyExec(Plan)**
 - Get GIDs of all nodes who has some specific value on a property. Returns a new result table.
 - Column-based storage, fast.
 - **SELECT gid FROM P_prop WHERE value = VALUE**
 - Slow on property tables that are not indexed

Execution Plans

- **FilterByConditionExec(ResultTable, Plan)**
 - Scan every item in result table, check if it has the value on the property to be checked. Drop those who does not have that value.
 - Supports inequations.

Execution Plans

- **FilterByConditionExec(ResultTable, Plan)**
 - Two execution methods:
 - A. If result table is small, or dataset is large:
Check every value in the previous result.
(Can be concurrently executed)
 - B. If dataset is small but result table is large:
Get all valid values from DB, and drop records with invalid values.
(Reduce slow DB lookups.)

Execution Plans

- **FilterRelationEqualityExec(ResultTable, Plan)**
 - Scan every item in result table, check if the two variables satisfies the equation in the plan. Keep those who satisfy.
 - One scan through the table.
 - No DB lookups. (Same EdgeID, same edge.)

Execution Plans

- **ExpandAllExec(ResultTable, Plan)**
 - Get all edges of that expanded variable in previous Result Table.
 - Put then into a set.
 - Reduce DB lookups.
 - Get all edges that starts from those values(nodes). Add those edges and nodes into result table. Drop those who owns no such edges.
 - If there are conditions on the end of that edge, keep only edges who satisfy these conditions.

Execution Plans

- **ExpandAllExec(ResultTable, Plan)**
 - Expand operation in ResultTable:
 - $(a)-[r_1] \rightarrow (b_1), (a)-[r_2] \rightarrow (b_2) \dots (a)-[r_n] \rightarrow (b_n)$
 - For each row, if value is a:
 - Concatenate all $[r_i] \rightarrow (b_i)$ to that row.
 - Create n new rows in result.
 - Need two scan through the whole table.
 - One to get the values to be expanded
 - One to add new values.

Execution Plans

- ExpandAllExec(ResultTable, Plan)

- Prev:

| A | R1 | B |
|-----|----|-----|
| n11 | r1 | n21 |
| n11 | r2 | n22 |
| n12 | r3 | n23 |
| n13 | r4 | n24 |

- Insert

| A | R2 | C |
|-----|----|-----|
| n11 | r5 | n31 |
| n11 | r6 | n32 |
| n13 | r7 | n33 |

Execution Plans

- ExpandAllExec(ResultTable, Plan)

- After:

| A | R1 | B | R2 | C |
|----------------|---------------|----------------|----|-----|
| n11 | r1 | n21 | r5 | n31 |
| n11 | r2 | n22 | r5 | n31 |
| n11 | r1 | n21 | r6 | n32 |
| n11 | r2 | n22 | r6 | n32 |
| n12 | r3 | n23 | | |
| n13 | r4 | n24 | r7 | n33 |

Execution Plans

- **ExpandIntoExec(ResultTable, Plan)**
 - If both ends of an edge are already in the result table, but we want the edge, we need ExpandInto.
 - Different execution methods considering the number of records in previous result.
 - Suppose expandInto operation is applied to **(a) – [r] ->(b)**, where **(a)** and **(b)** are already found, and we need to expand **[r]**.

Execution Plans

- **ExpandIntoExec(ResultTable, Plan)**
 - A. If table is really small:
 - Check every a-b pair, get all edges from a to b.
 - Similar to ExpandAll.
 - **Small** : the execution time of performing #Small checks does not exceed the time to get all values in the edge table.
(time to execute 3000 single query = time to query the whole table with 2M records)

Execution Plans

- **ExpandIntoExec(ResultTable, Plan)**
 - B. If table is larger, but not as large as the Edge table:
 - Get possible (a)-(b) pairs.
 - Put them into a set.
 - Find all edges by these pairs.
 - Edge table supports edge lookup by combining multiple conditions.
 - Rest of work is similar to ExpandAll
 - Much faster if many duplicated (a)-(b) pairs

Execution Plans

- **ExpandIntoExec(ResultTable, Plan)**
 - C. If table is larger than the total number of edges in dataset:
 - Get all edges from DB.
 - Construct (a)-(b) pairs of every record in the result table.
 - Get all edges starts from a and ends at b.
 - In-memory lookup. Can be faster by creating index on result table.
 - Cost of scanning result table is more than cost of getting all edges.

Execution Plans

- **CartesianProductExec(Table1, Table2, Plan)**
 - Cartesian product on two tables. Two tables have no common variables.
 - Nested for-loops.
 - Result is stored in table1, drop table2.
 - Saves memory space

Execution Plans

- **NodeHashJoinExec(Table1, Table2, Plan)**
 - Join on two tables with same variable or with equation between two tables.
 - Hash is calculated by multiplication of `hashcode()` of all the join variables(common variables)
 - Hashed variables are GID of a node (as a string), or the string value of a Property Lookup.

Execution Plans

- **NodeHashJoinExec(Table1, Table2, Plan)**
 - Steps:
 - Calculate hash value of all rows in two tables.
 - Use a map to store the hash values of table1:
 - Key : Hash
 - Value: indices of records with this hash to the table1
 - Java Hashmap provides $O(1)$ lookup.

Execution Plans

- **NodeHashJoinExec(Table1, Table2, Plan)**
 - Steps:
 - For each row in table2, get the hash value of that row
 - Lookup the hash in table1.
 - Perform Cartesian Product on rows with same hash. (If there are multiple rows with same hash value). Insert result into new table
 - Return new table.

Execution Plans

- **NodeHashJoinExec(Table1, Table2, Plan)**

- Result Table implementation:

```
new_table = []
foreach row : table2
    hash = hash_func(row)
    if hash in table1.hashmap:
        indices = table1.hashmap.get(hash)
        table1_rows = table1.get_rows(indices)
        new_rows =
            cartesian_product(row,
table1_rows)
        new_table.add(new_rows)
```

Execution Plans

- **RangeExpandAllExec(Table1, Plan, PathList)**
 - Looks for all possible variable length path expansions on the start node.
 - Get all conditions on edge variable and end node variable.
 - Get the starting nodes, and BFS on each of them searching for path, tracking passed nodes and edges.

Execution Plans

- **RangeExpandAllExec(Table1, Table2, Plan, PathList)**

```
queue <- []
result <- []
for i <- (0 to range.end):
    for each path in queue:
        if path length is within range:
            if path ends at a valid node:
                result <- path

// BFS
new_queue = []
for path in queue:
    edges <- getNeighbors(path.end_node)
    for edge in edges:
        new_path <- path + edge + edge.end
        new_queue <- new_path
queue = new_queue
```


Execution Plans

- **RangeExpandAllExec(Table1, Plan, PathList)**
 - Result is stored in pathlist, and kept globally.
 - Path is represented by index to itself in the pathlist.
 - Cons: Saves space. Faster.
 - Pros: Path list is append only. No modifications.

Execution Plans

- **RangeExpandIntoExec(Table1, Plan, Pathlist)**
 - Similar to RangeExpandAll.
 - Difference is that end node should be in same as the end node in the previous result.

Execution Plans

- **ProduceResultExec(Table1, Plan, Pathlist)**
 - Produce the final result of the query.
 - A separated List<Object> to store final results.
 - Expand intermediate result by its IDs.
 - Node: ExpandNode() to get all contents of a node by its GID.
 - Edge: ExpandEdge() to get edge Eid, start node GID and end node GID.
 - PropertyLookup: Just get the string
 - Path: Fetch the path by index to Pathlist.

Working example

Example

- Query String:

match

(a)-[r]->(b),

(c)-[r2]-(b),

(a)--(c)

where

b.id = "8"

return

a.id, r, b.name

Example

- **Query String:**

match

(a)-[r]->(b),

(c)-[r2]-(b),

(a)--(c)

where

b.id = "8"

return

a.id, r, b.name

Query Plan:

ScanById|b|

ExpandAll|b|-[r]-(a)

ExpandAll|b|-[r2]-(c)

FilterRelationEquality||r2 != r

ExpandInto|a|-[anonRelation0]-(c)

FilterRelationEquality||anonRelation0
!= r AND anonRelation0 != r2

ProduceResult|a,b|a:a.id b:b.name

Example

- **Query Result:**

...(many rows)

| | | |
|---|----------------------------------------------------------------------------------------------------------|--------------------------|
| | ===== | |
| 0 | {node2Label=4, eid=2, node2=2, rel_type=HashTagUsedBy, node1Label=3, relationship=null, node1=1} | ThinkBIGSundayWithMarsha |
| 0 | {node2Label=4, eid=2, node2=2, rel_type=HashTagUsedBy, node1Label=3, relationship=null, node1=1} | ThinkBIGSundayWithMarsha |
| 2 | {node2Label=4, eid=276, node2=2, rel_type=HashTagUsedIn, node1Label=2, relationship=null, node1=261} | ThinkBIGSundayWithMarsha |
| 2 | {node2Label=4, eid=276, node2=2, rel_type=HashTagUsedIn, node1Label=2, relationship=null, node1=261} | ThinkBIGSundayWithMarsha |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel_type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |
| 4 | {node2Label=4, eid=29270, node2=2, rel type=hashTagComenation, node1Label=4, relationship=null, node1=3} | ThinkBIGSundayWithM |

Example

- SQL Statements

...(many rows)

SQL:

=====

```
SELECT gid FROM P_id WHERE value = "8";
SELECT node1, eid FROM Edge WHERE node2 = "2" ;
SELECT node1, eid FROM Edge WHERE node2 = "2" ;
SELECT node2, eid FROM Edge WHERE node1 = "2" ;
SELECT node2, eid FROM Edge WHERE node1 = "1" ;
SELECT node2, eid FROM Edge WHERE node1 = "3" ;
SELECT node2, eid FROM Edge WHERE node1 = "261" ;
SELECT node1, eid FROM Edge WHERE node1 = "261" AND node2 = "1" ;
SELECT node1, eid FROM Edge WHERE node1 = "261" AND node2 = "3" ;
SELECT node1, eid FROM Edge WHERE node1 = "261" AND node2 = "261" ;
SELECT value FROM P_id WHERE gid = "1";
SELECT * FROM edge WHERE eid = "2";
SELECT value FROM P_name WHERE gid = "2";
SELECT value FROM P_id WHERE gid = "1";
SELECT * FROM edge WHERE eid = "2";
SELECT value FROM P_name WHERE gid = "2";
SELECT value FROM P_id WHERE gid = "261";
SELECT * FROM edge WHERE eid = "276";
SELECT value FROM P_name WHERE eid = "2".
```


Example

- **Query String:**

```
match (a)-[r]->(b), (c)-[r2]-(b), (a)--(c) where  
b.id = "8"
```

```
return a.id, r, b.name
```

Future work

Future work

- **DB Lookup is expensive**
 - Some data can reside in memory.
 - Perform batch lookup, and use call-backs.
 - Add cache on Database Utility
- **Concurrency improves a lot.**
 - Concurrent execution on the execution plans.
 - Plans of same depth in the planning tree can be executed concurrently.
 - Concurrent execution on neighbor lookup in variable length expansion

Future work

- **Nodes are biased**
- (Kurant M, Markopoulou A, Thiran P. On the bias of BFS (breadth first search)[C]//Teletraffic Congress (ITC), 2010 22nd International. IEEE, 2010: 1-8.)
 - Most of nodes have low degree. But some have really high degree.
 - BFS is not efficient.
 - Apply parallel BFS algorithm.
- **Partitioning on database**
 - Connected graph
 - On biased nodes

Future work

- **Dataset is huge**
 - Lazy evaluation on expensive SQL queries
 - Column-based intermediate result table
 - Partitioning also helps.
- **Execution Plan optimization**
 - Reordering
 - Combining related plans
 - Intermediate result sharing

Thanks